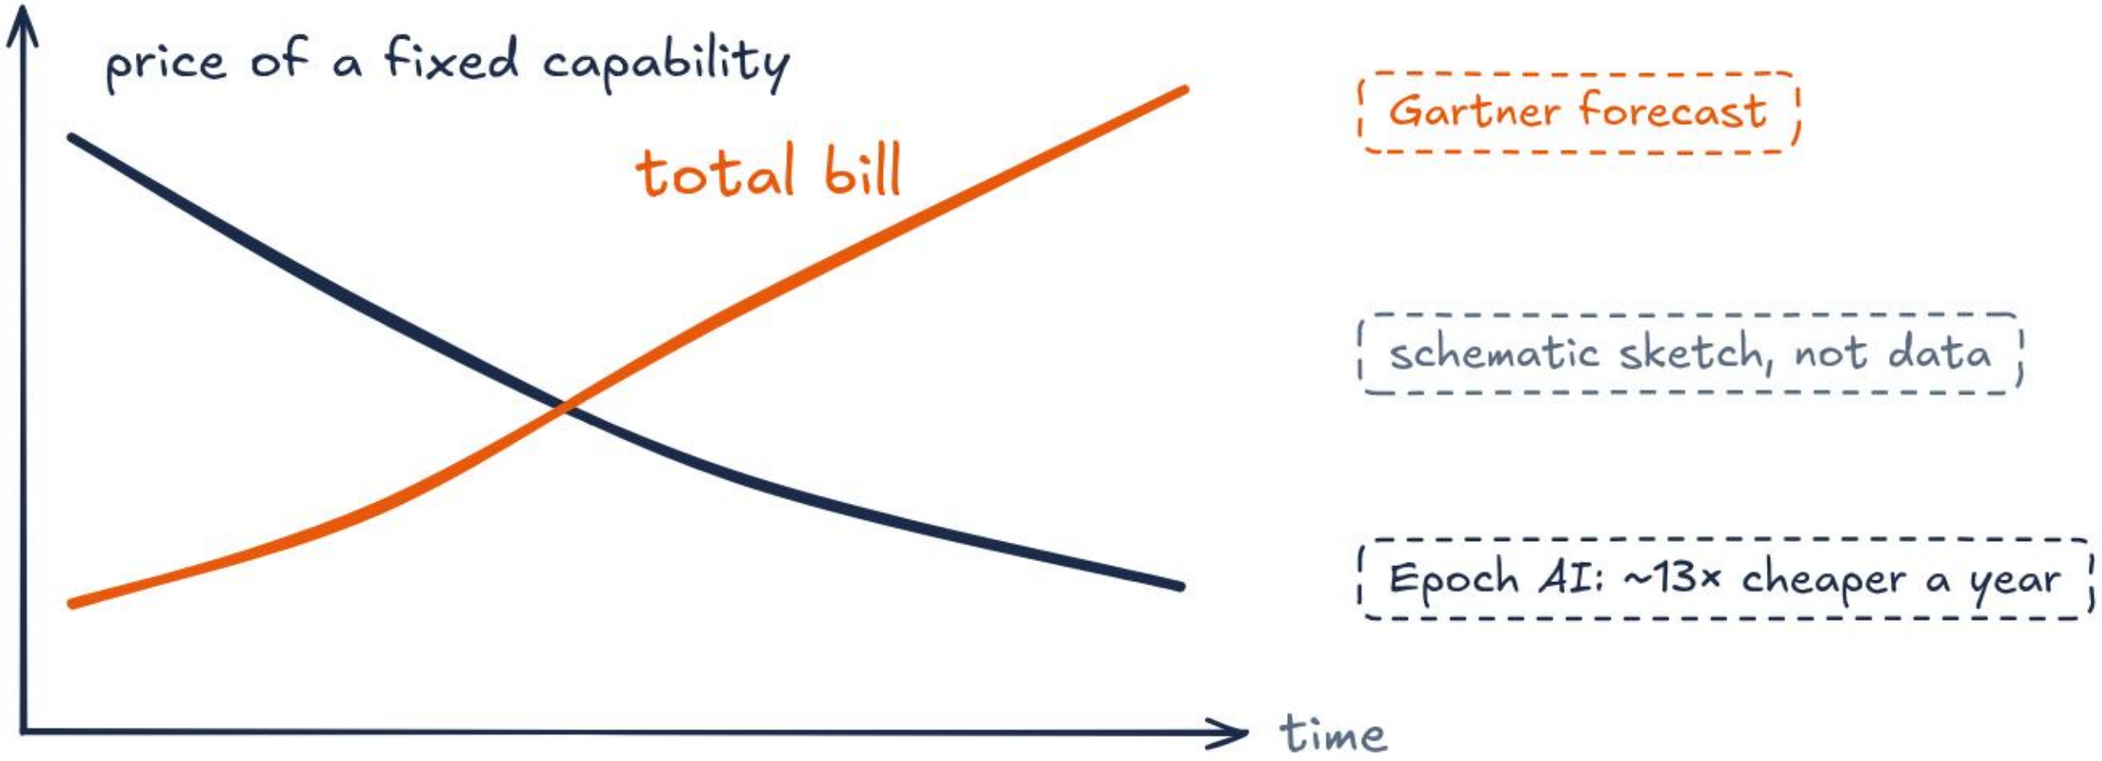
Internal talk · 30 min

# Taming the AI bill

## Cost control for AI features

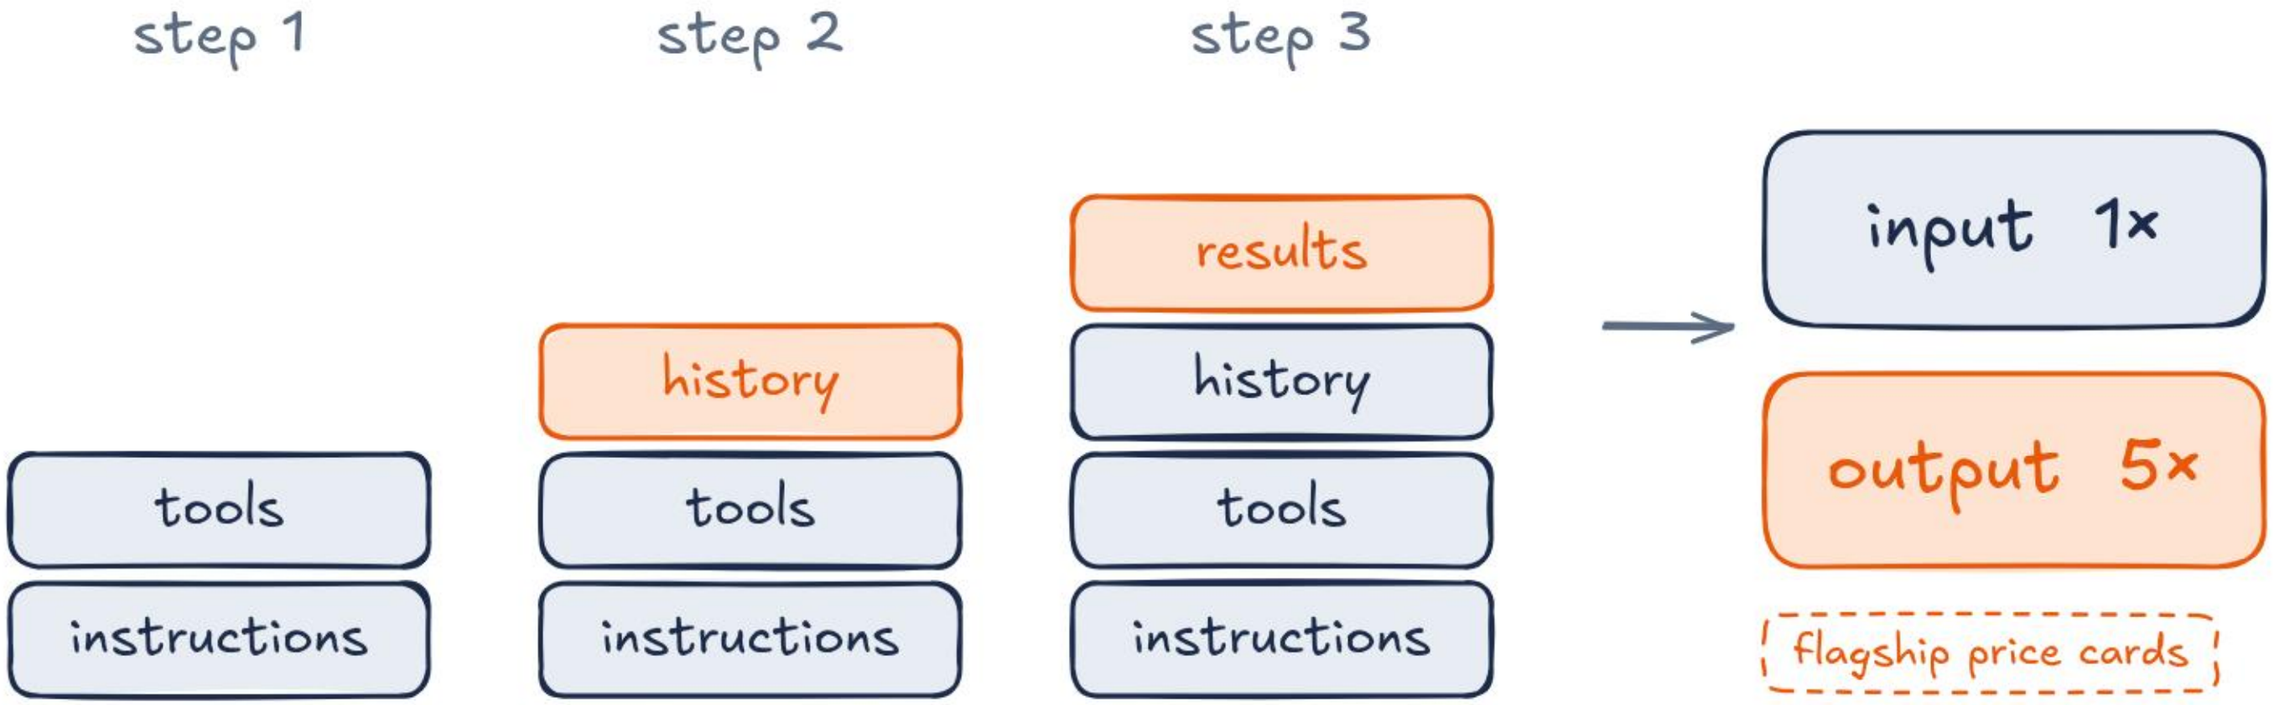


# Cheaper tokens, bigger bills




Unit price falls. Spend rises.

# Every step re-sends everything




cost grows with steps × context

# Easy-to-miss multipliers

| ×2–5 | ×2–6 | +10% | ×2 |
| --- | --- | --- | --- |
| **Long prompts** | **Fast tier** | **Data region** | **Promo ends** |
| above a size tier | pay for speed | residency pricing | one model, 1 Dec |

as read on 19 Oct 2026 · check before quoting

Batch and Flex tiers: -50%

# Cost per task, not per token




example: 14 of 16 tickets pass, 100k tickets a month

our measurement · frontier price card

a cheap wrong answer is not cheap

# FinOps for AI in three verbs



## Measure

per use case




## Allocate

per team and feature



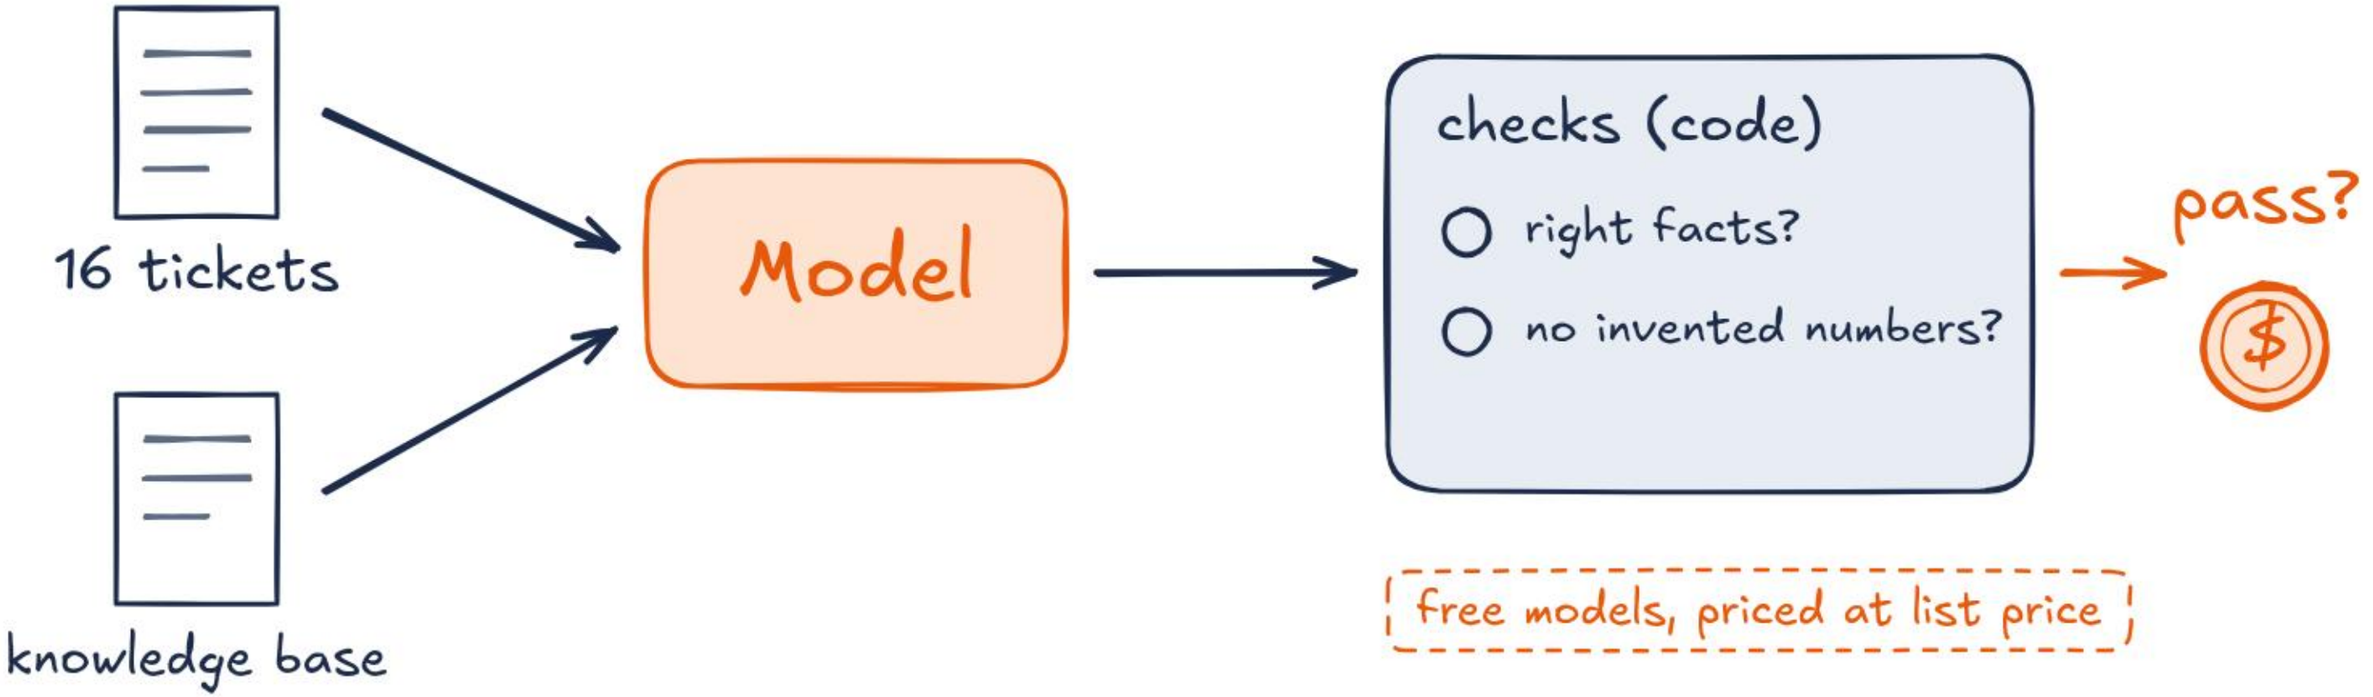
## Limit

budgets and alerts

**98%** of teams now manage AI spend

31% two years ago · State of FinOps 2026

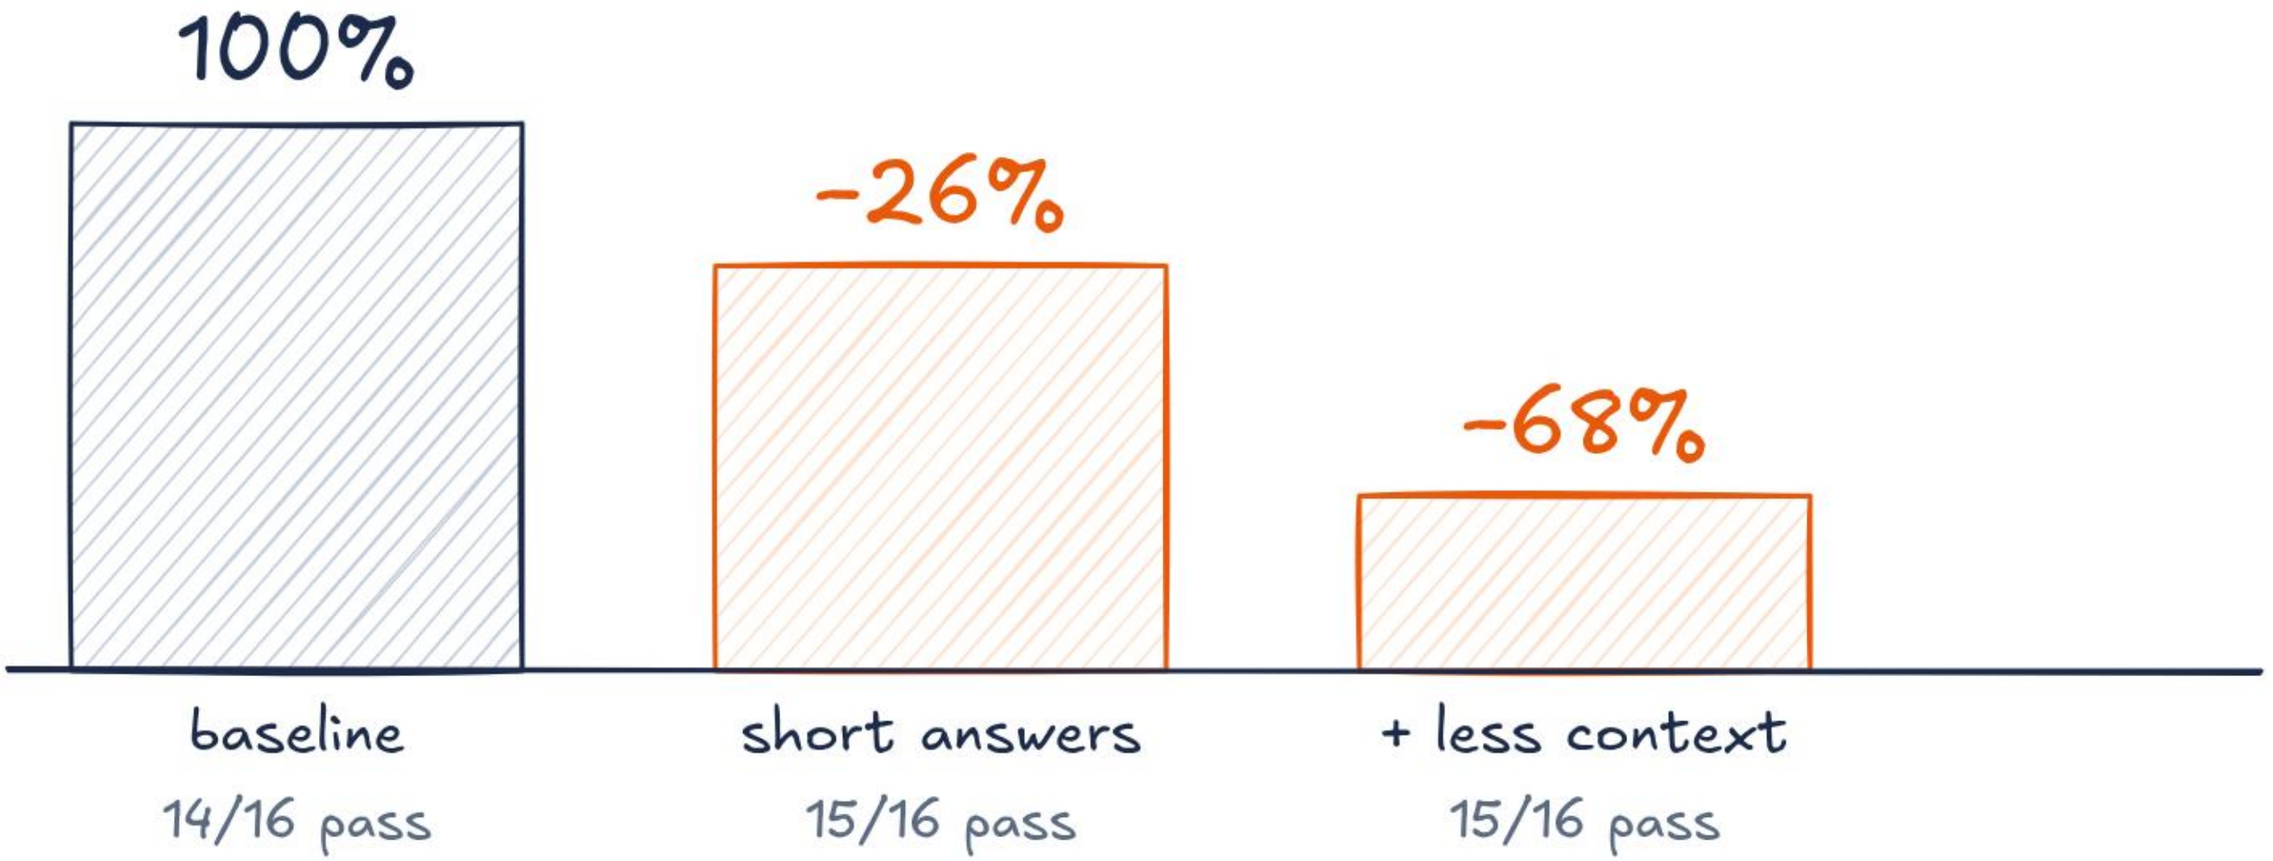
# Our test bench




same tickets, one lever at a time

# Say less, send less

our measurement



# A smaller model

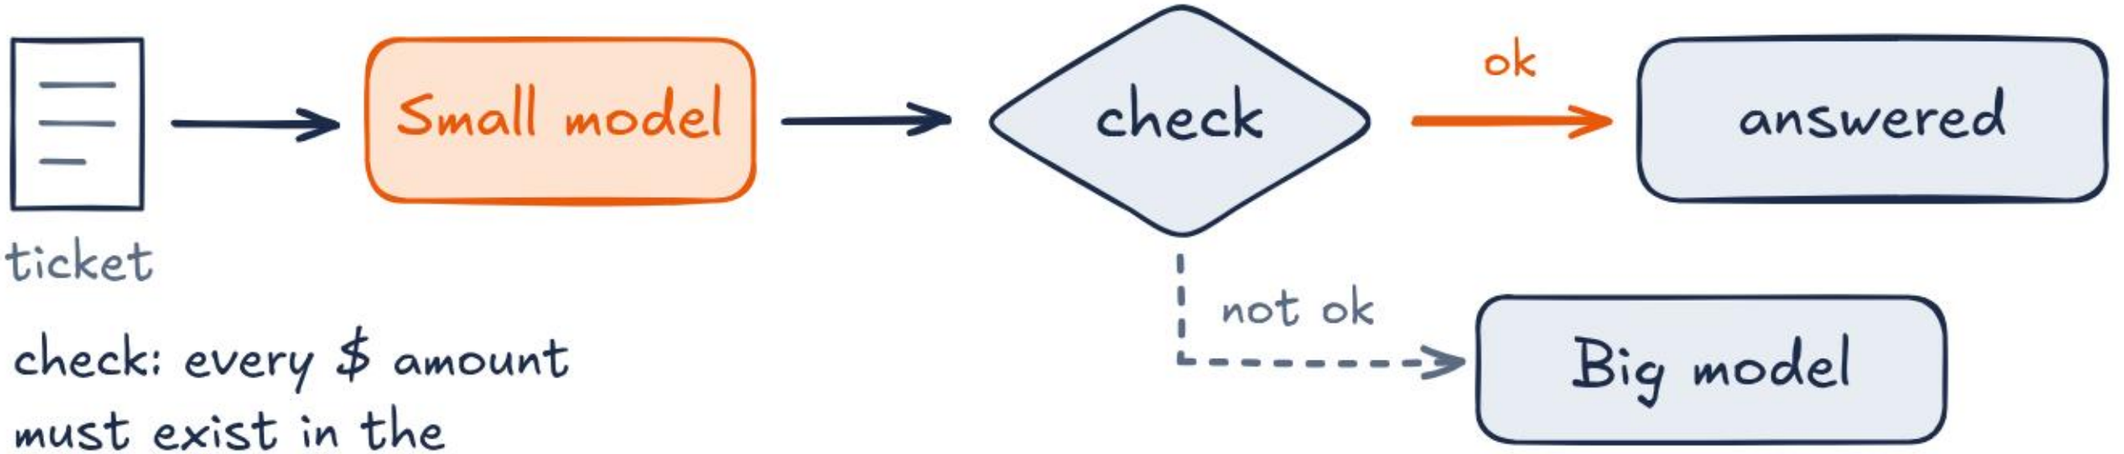
our measurement

| big model | small model | small + less context |
| --- | --- | --- |
| 100% | -96% | -98% |
| cost | cost | cost |
|  |  |  |
| 14/16 pass | 14/16 pass | 9/16 pass |

levers interact: re-test after each one

# A cascade needs a checker

our measurement




check: every $ amount must exist in the knowledge base

✗ self-reported confidence
13/16

✓ number check
14/16 at -95%

# Prompt caching: keep the start stable

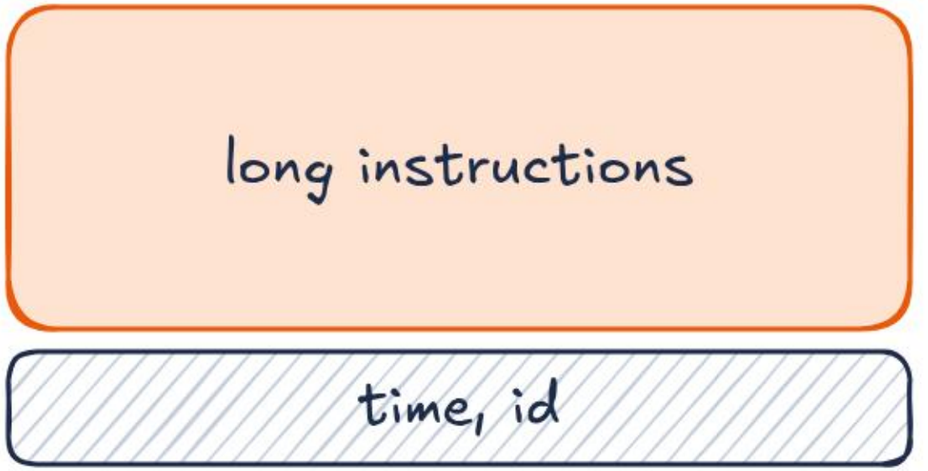
## dynamic part first




**0**

cached tokens per request

## dynamic part last



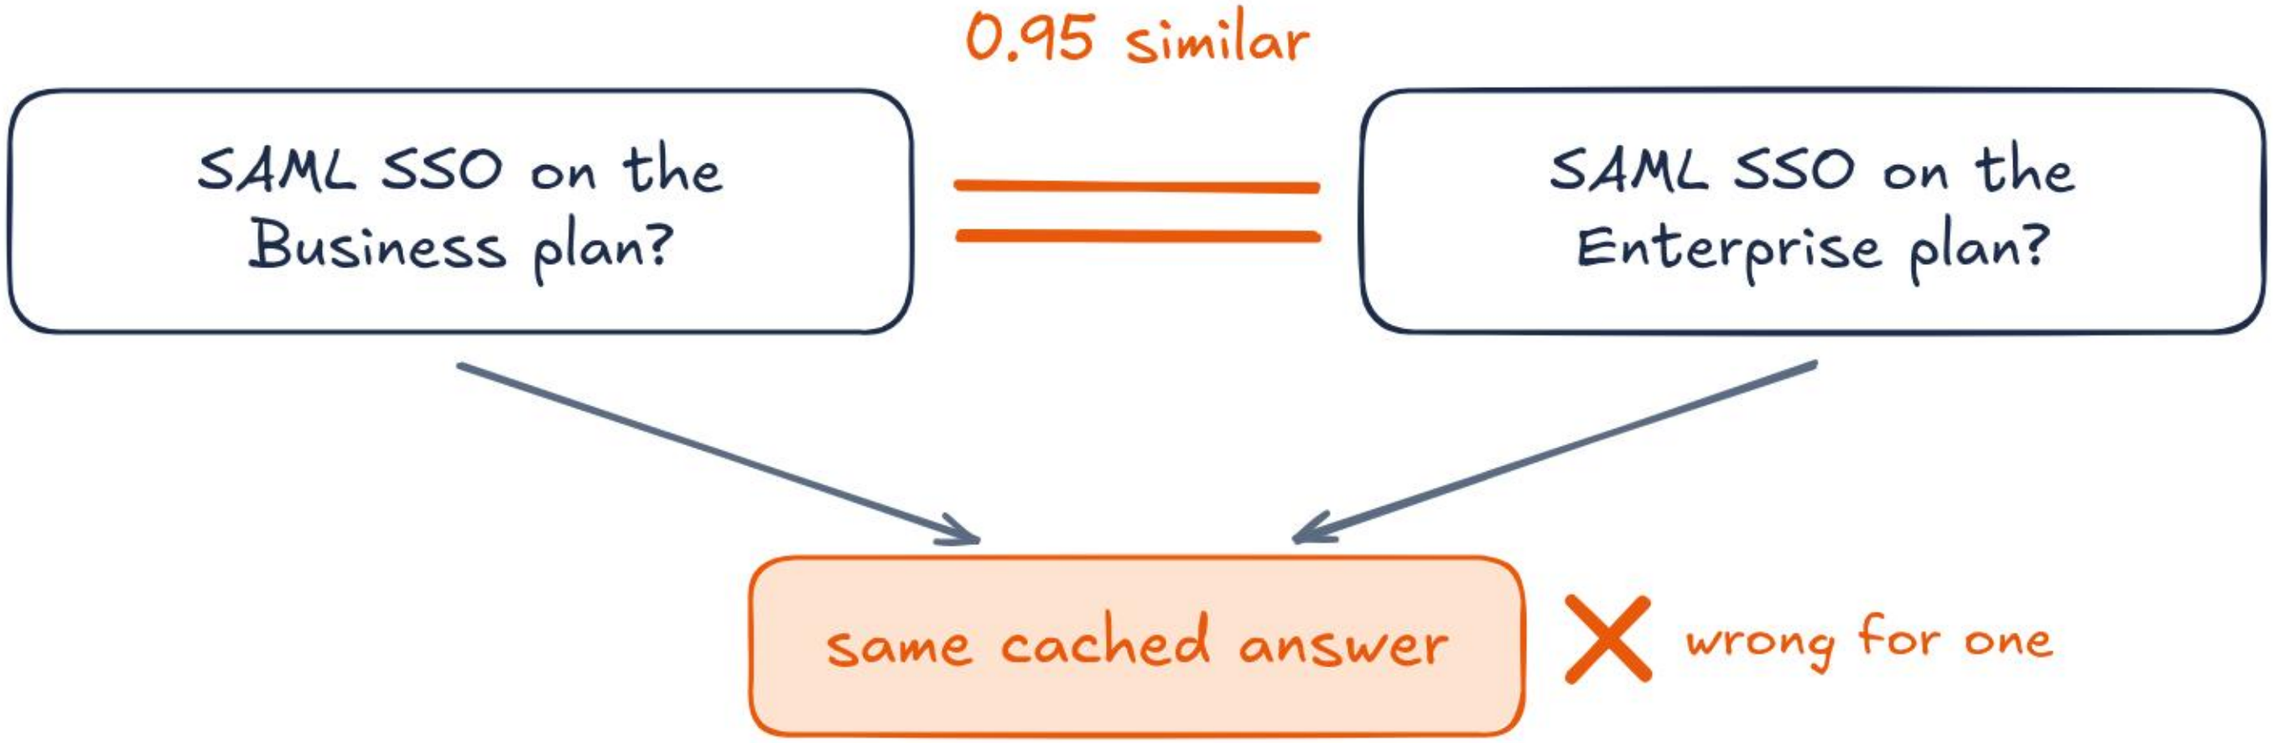

**6,288**

cached tokens per request

our measurement · ~8k-token prompt · free endpoint

paid rules: reads ~0.1×, writes ~1.25×

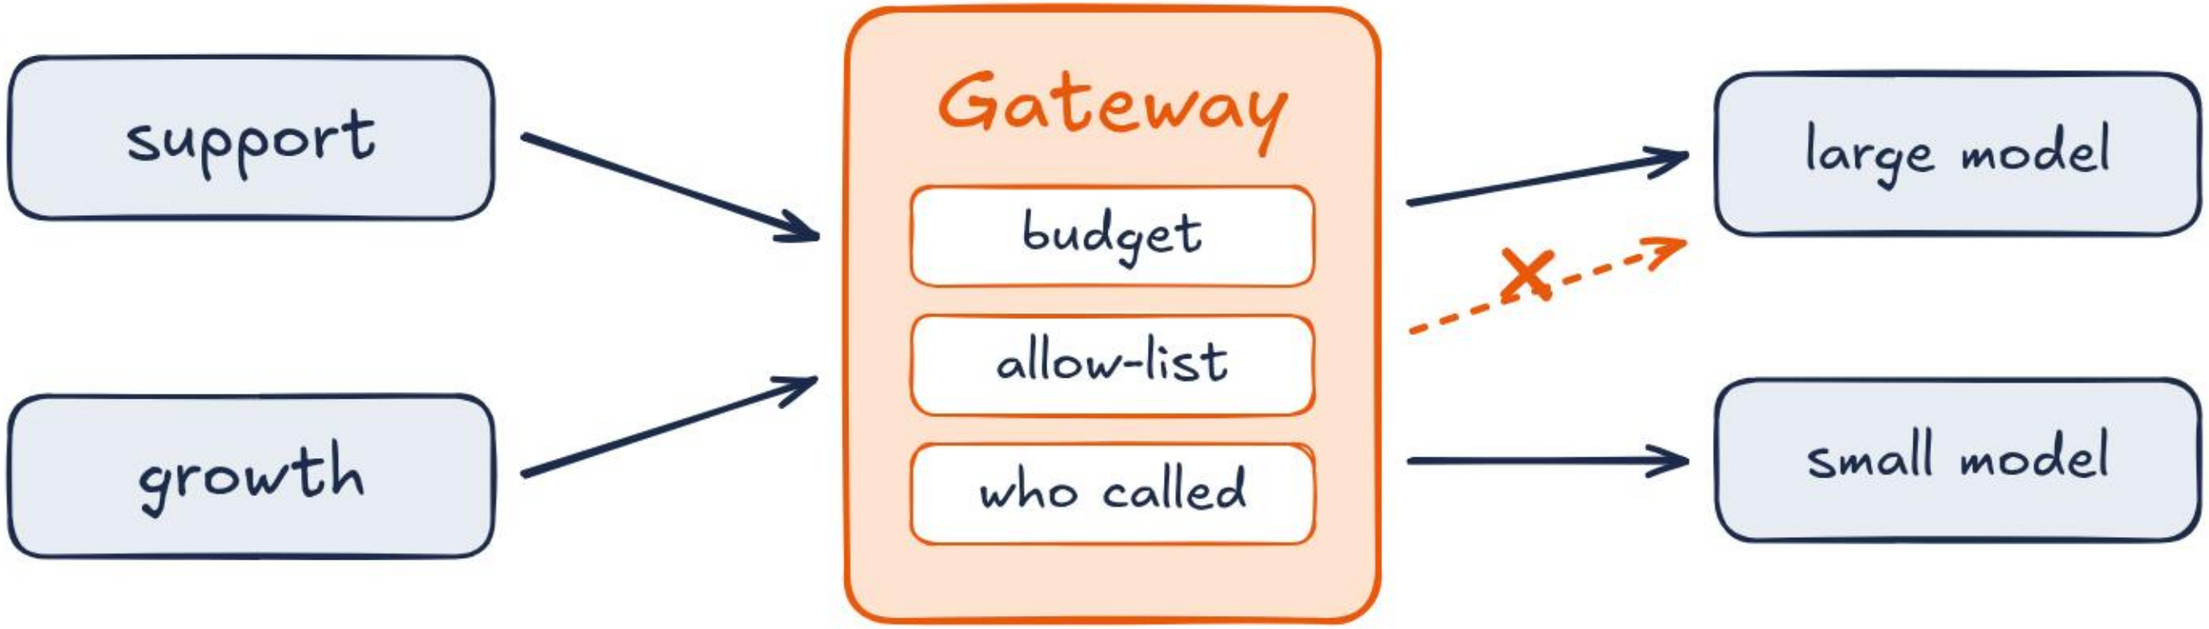
# Similar is not the same




our stress test · 24 pairs built to be tricky

a wrong answer is free and instant

# One door for every call




support: stopped at 112% of budget
growth: large model refused

our measurement

one request of overshoot is normal

# Live demo: the cost table

```
config          pass   $/task   vs base
A_baseline      14/16 $0.01250      +0%
B_concise       15/16 $0.00928     -26%
D_trim          15/16 $0.00394     -68%
F_small_full    14/16 $0.00046     -96%
G_small_trim     9/16 $0.00019     -98%
```

→ cheaper is not always better

$ python ex2_levers/run_levers.py --rescore results/...

saved run · shortened table · frontier price card

needs no key: it re-scores our committed runs

# Live demo: the $0 trap

```
small:
  model: openai/nvidia/…:free

small-openrouter-prefix:
  model: openrouter/nvidia/…:free
```



```
small
  cost logged: 2.8e-06

small-openrouter-prefix
  cost logged: 0.0
```

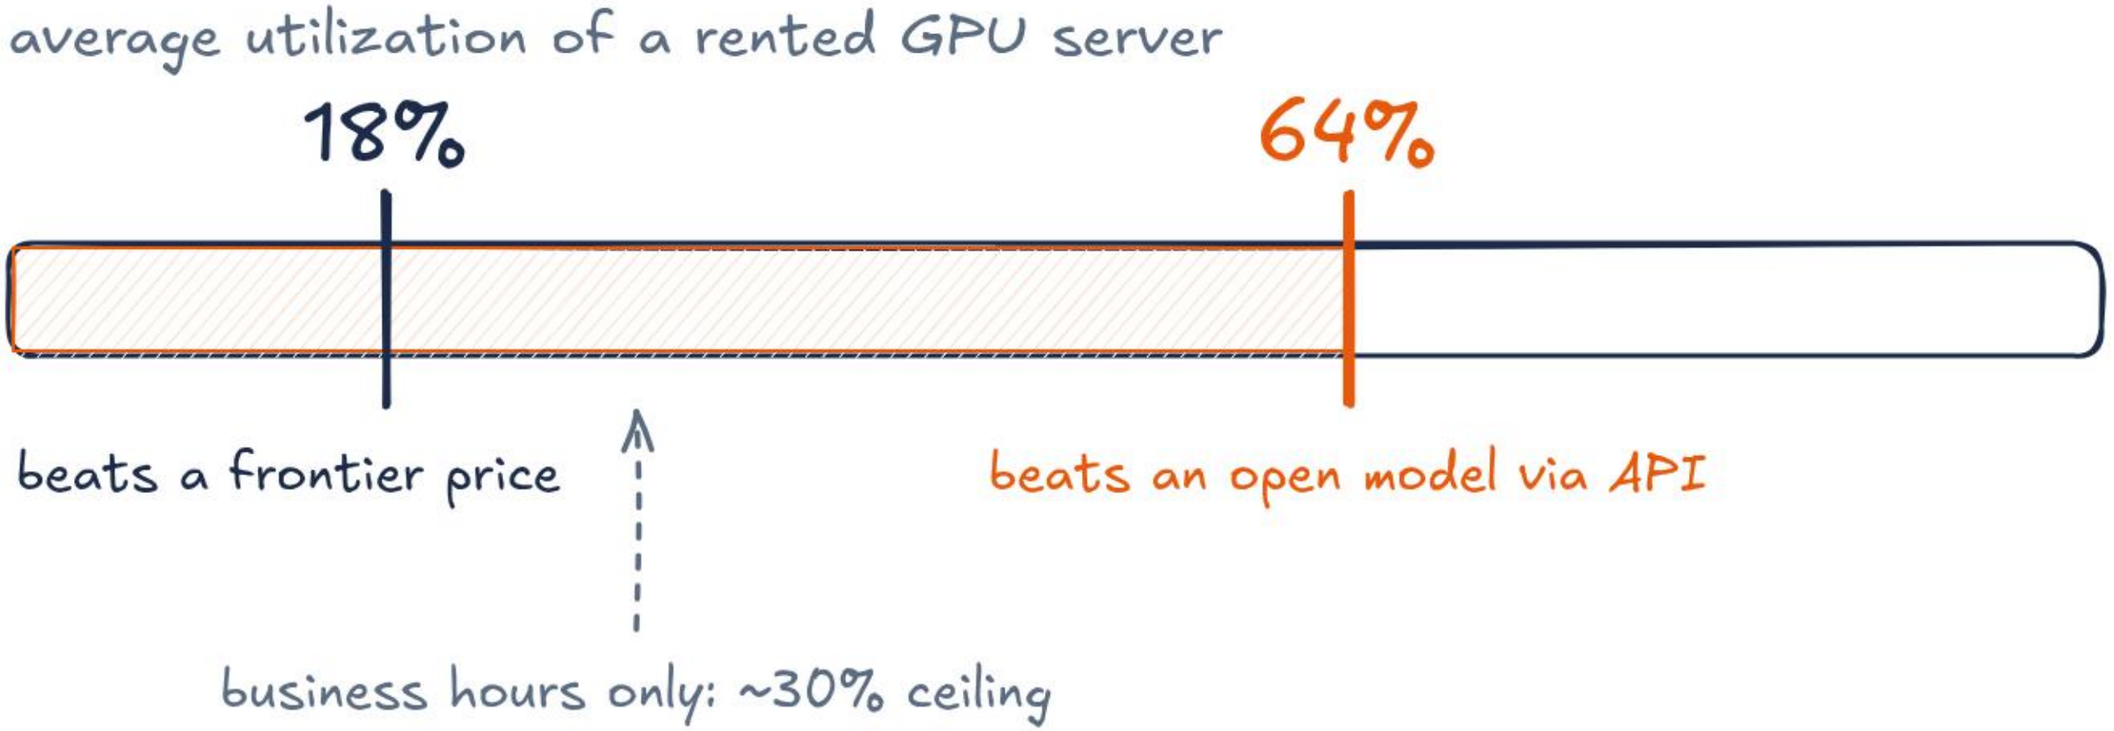
same model, same custom price

18 in × $0.10/M + 2 out × $0.50/M = $0.0000028

check one request against a hand calculation

saved run · LiteLLM v1.106.0 · free model, shadow prices

# When does self-hosting pay?




our arithmetic · throughput is an assumption

# Honest limits

**16 tickets, one run**
1–2 tickets is noise

**Free models, list prices**
shadow cost, nothing charged

**Caching: partly simulated**
paid rules applied by math

**Survey numbers**
reported by their publishers

# Three takeaways

1. Count cost per successful task
2. One lever at a time, re-test quality
3. One gateway door, check its numbers

# Try it yourself

https://github.com/vusatui/rnd-research
→ ai-finops/

- Run the workshop
- Re-score our runs
- Read the research
- Compare the gateways



scan me

# Questions?

Taming the AI bill